

Welcome to Beginner Track!

Workshop 1: Intro to Machine Learning

Attendance code: **Amazon**

Linktree: bit.ly/btrack-w22-feedback

Discord: bit.ly/ACMdiscord

Our Mission

To build and develop a community of students interested in **Artificial Intelligence** at UCLA and beyond.

Our Values

- Technical Proficiency and Awareness in Artificial Intelligence
- Creating a Positive Impact on Society
- Diversity and Inclusion

ACM AI Initiatives

Workshops

Events



acm.ai

Outreach

Projects

Projects

Get a chance to work on your first deep learning project!

Last quarter, our teams built deep learning models to:

- Classify plant diseases
- Identify humpback whales

Only prerequisite is Advanced Track!

Links: [FAQ](#), [Interest form](#) 📌



Our Workshops



- **Beginner Track – *What is ML?***

- Basics of machine learning
- Implement linear and logistic regression



- **Advanced Track – *Deep Learning***

- Concepts like deep neural networks, CNNs, RNNs
- Basic knowledge of ML concepts expected

Beginner Track

Who's it for?

- no experience in machine learning
- minimal experience coding
- want a solid foundation in the theory behind ML

What's covered?

- basics of machine learning
- theory and implementation of simple models
- introduction to useful ML libraries

When and where are meetings?

- **Time:** Thursdays 6–8pm (Zoom)



Sudhanshu
Agrawal (he/him)



Jenson Choi
(he/him)



Claire
Huang



Naman Modani
(he/him)

Beginner Track: Schedule


- Workshop 1 (1/13): Intro to ML + Intro to Python
- Workshop 2 (1/20): K-Nearest Neighbours
- Workshop 3 (1/27): Linear Regression
- Workshop 4 (2/3): Logistic Regression
- Workshop 5 (2/10): Multiclass Classification
- Workshop 6 (2/17): Numpy and Pandas
- Workshop 7 (2/24): Guided Project
- Workshop 8 (3/3): Guided Project (continued)

Don't worry!

- Machine Learning can be daunting!
- We've got you! We'll walk you through all details and try to get you as comfortable with the math and coding sections as possible

AI and ML in Real Life

Computer vision



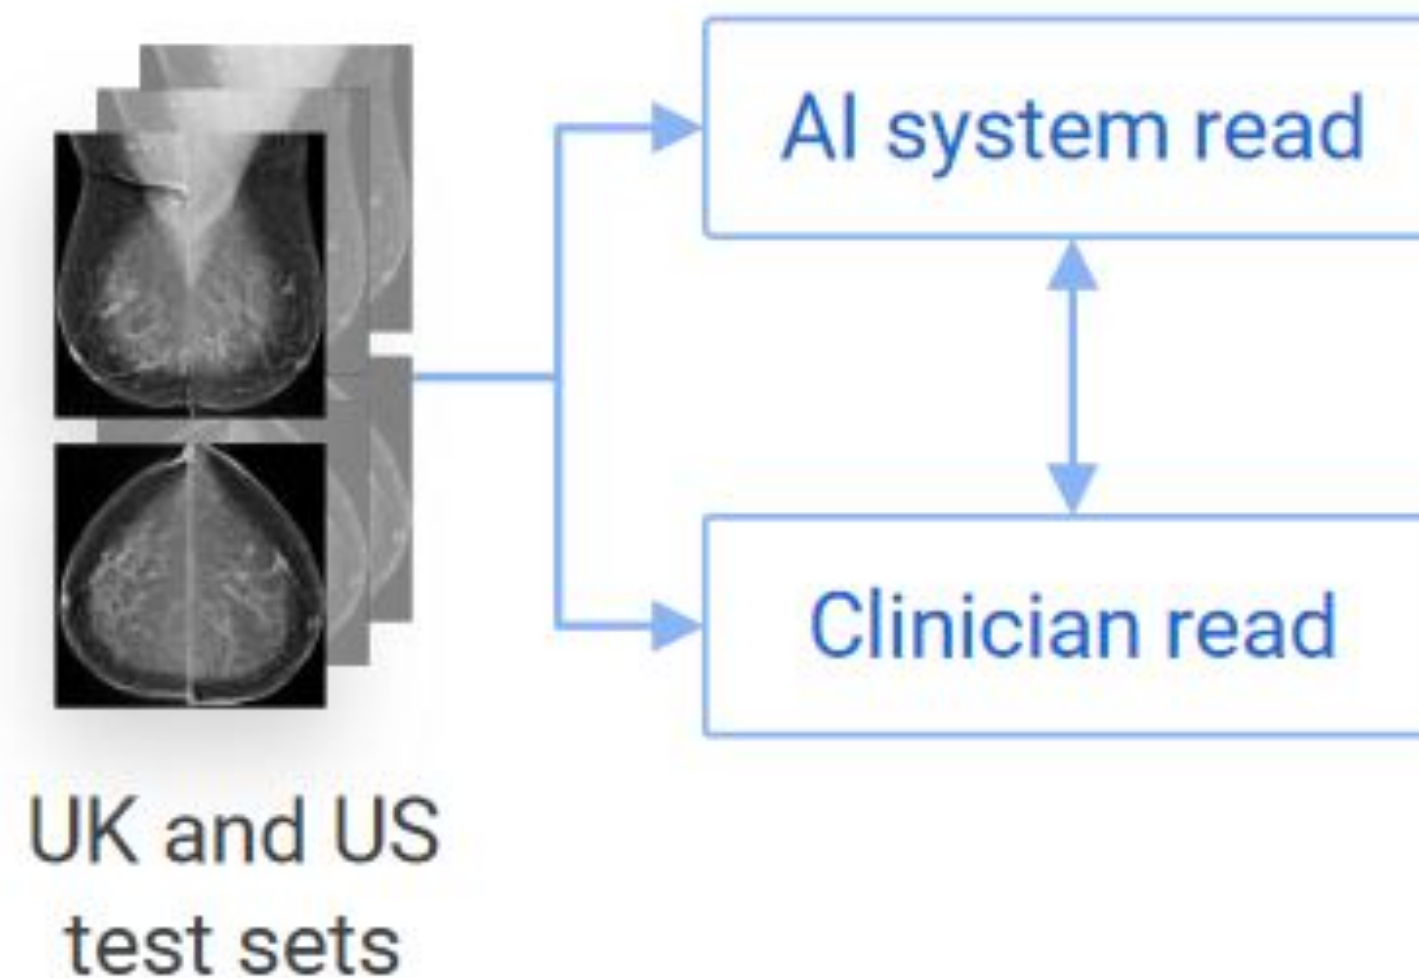
Convolutional neural networks have achieved stunning results in computer vision!

Source: [*Taking It to the Streets: Ride in an NVIDIA Self-Driving Car with DRIVE Labs*](#)

Healthcare

Evaluation

Comparison with retrospective clinical performance



Deep Learning techniques **outperform** trained specialists in some medical recognition tasks.

Image from: [International evaluation of an AI system for breast cancer screening](#)

BBC Article: <https://www.bbc.com/news/health-50857759>

Natural language processing



- <https://play.aidungeon.io/>
- Built with OpenAI's GPT-3 model
- Type anything you want!



OpenAI, Public domain, via Wikimedia Commons

The Intuition Behind ML


**Let's play a
game**



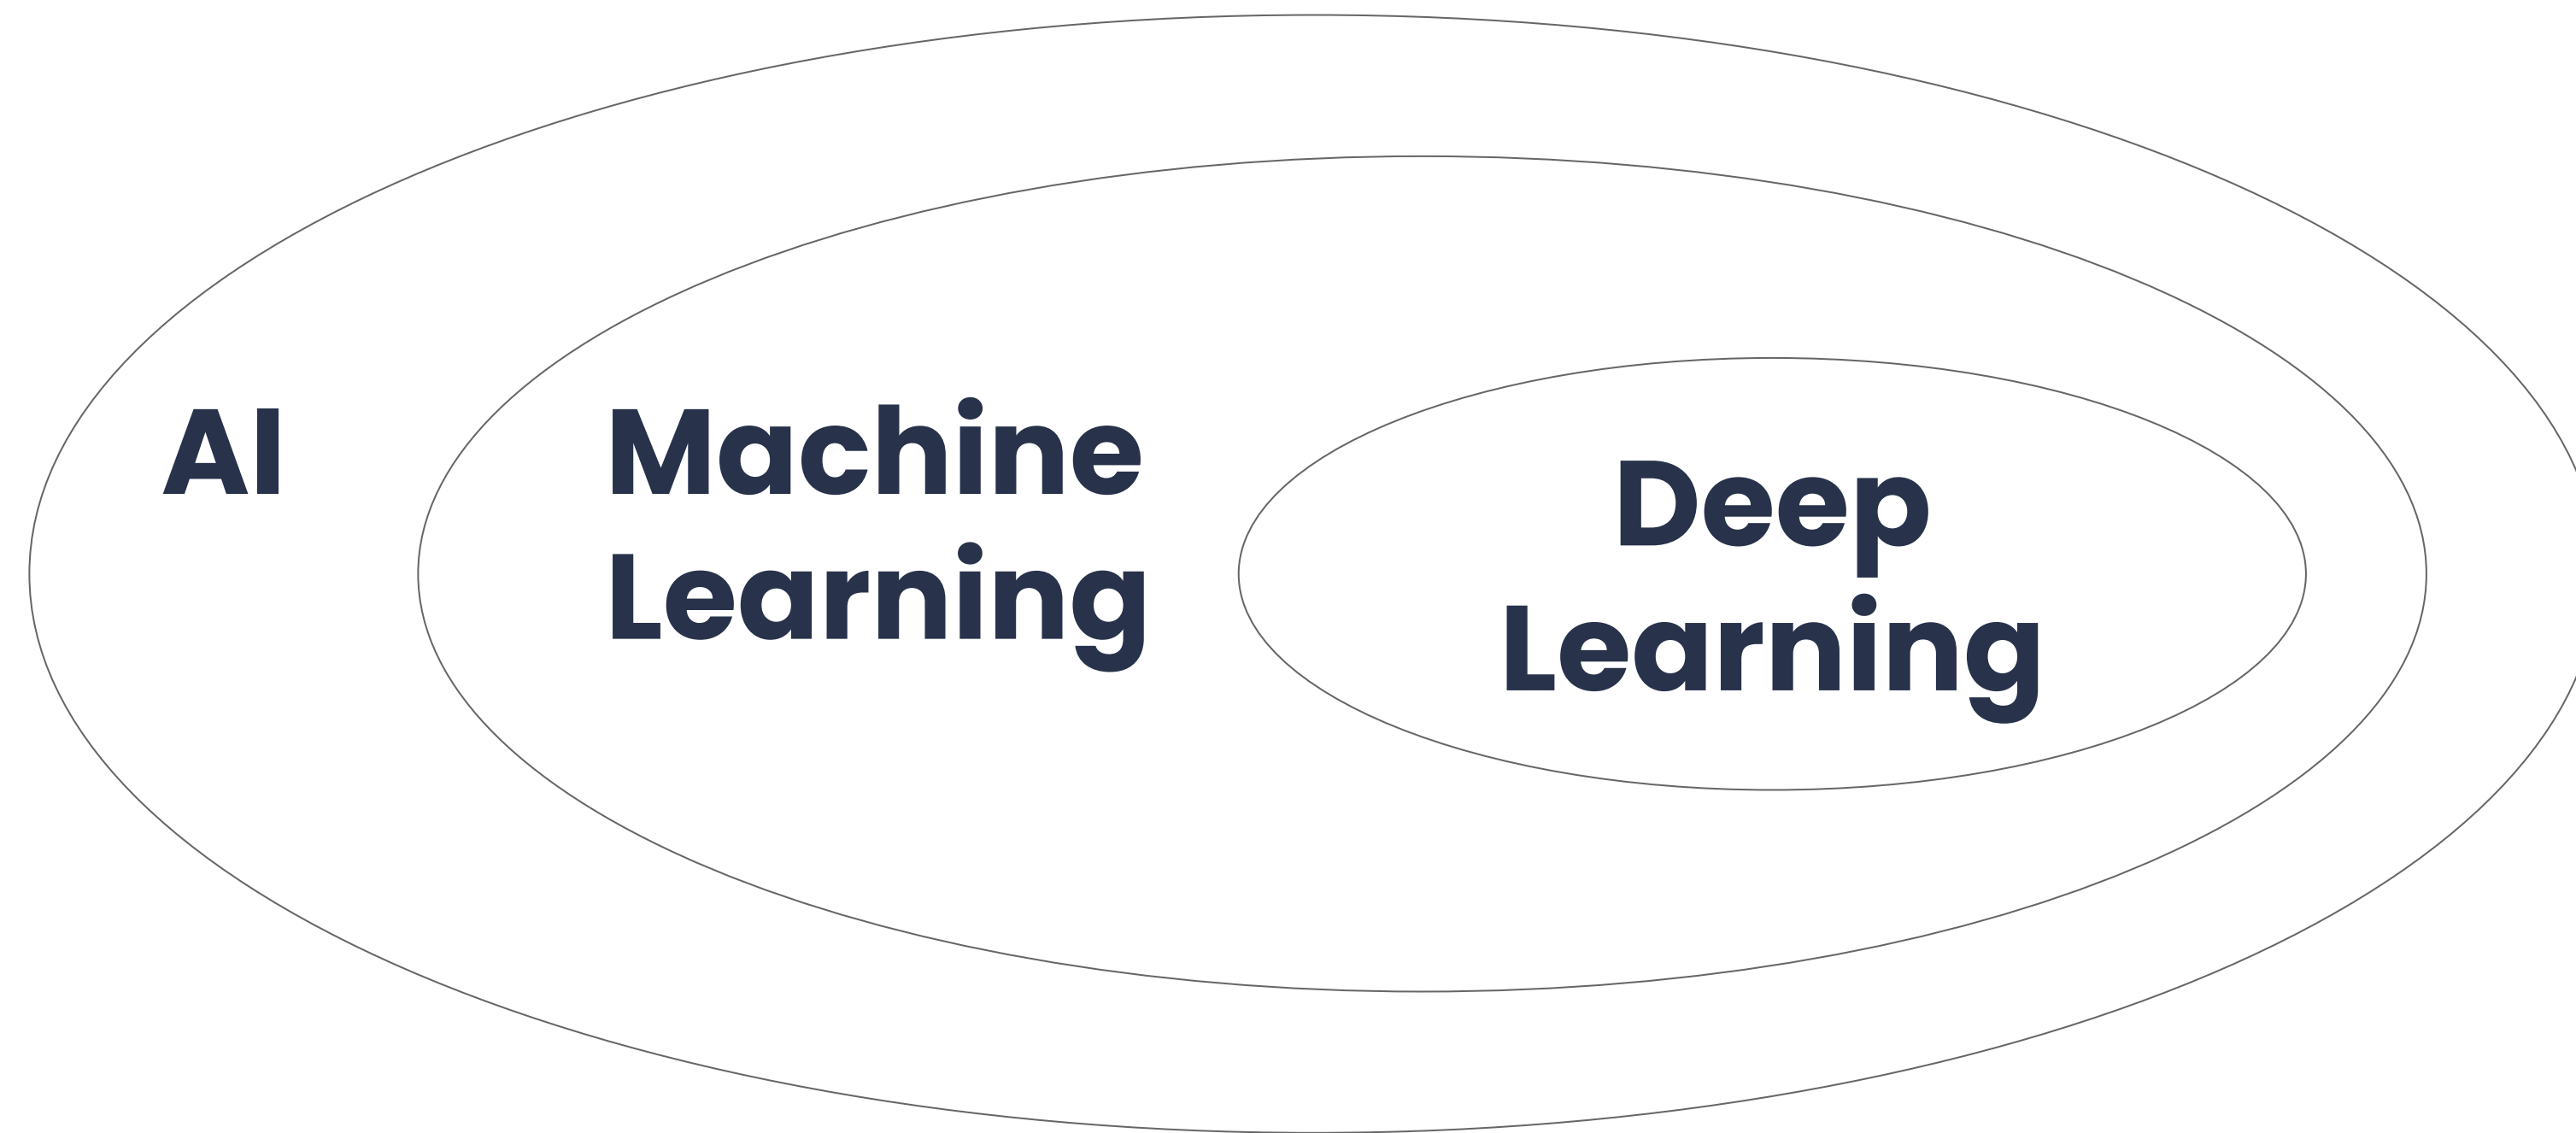
50-50

- The objective of the game is to find the **letter of the alphabet** such that **50%** of the audience's **names** come **before** this letter, and 50% **after**
- We're going to start off by choosing a random letter
- Each round we must choose one of the 3 options: [The first letter of my name is] Before the current letter, After the current letter, At the current letter

What is ML?



AI vs ML vs Deep Learning



Definitions

Artificial Intelligence – A concept

- Allowing computers to perform tasks that normally require human intelligence
- Eg. seeing, hearing, moving, decision making

Machine Learning – *A type of AI*

- A set of methods which can be used to allow computers to perform AI tasks without being explicitly programmed to do so.

Let's Discuss

Suppose we wanted to predict the price of a dorm room.
How would we go about doing it?



What is a model?

- “Something” that takes in an input and produces an output

Eg. Takes in a picture and determines whether it is a cat or dog

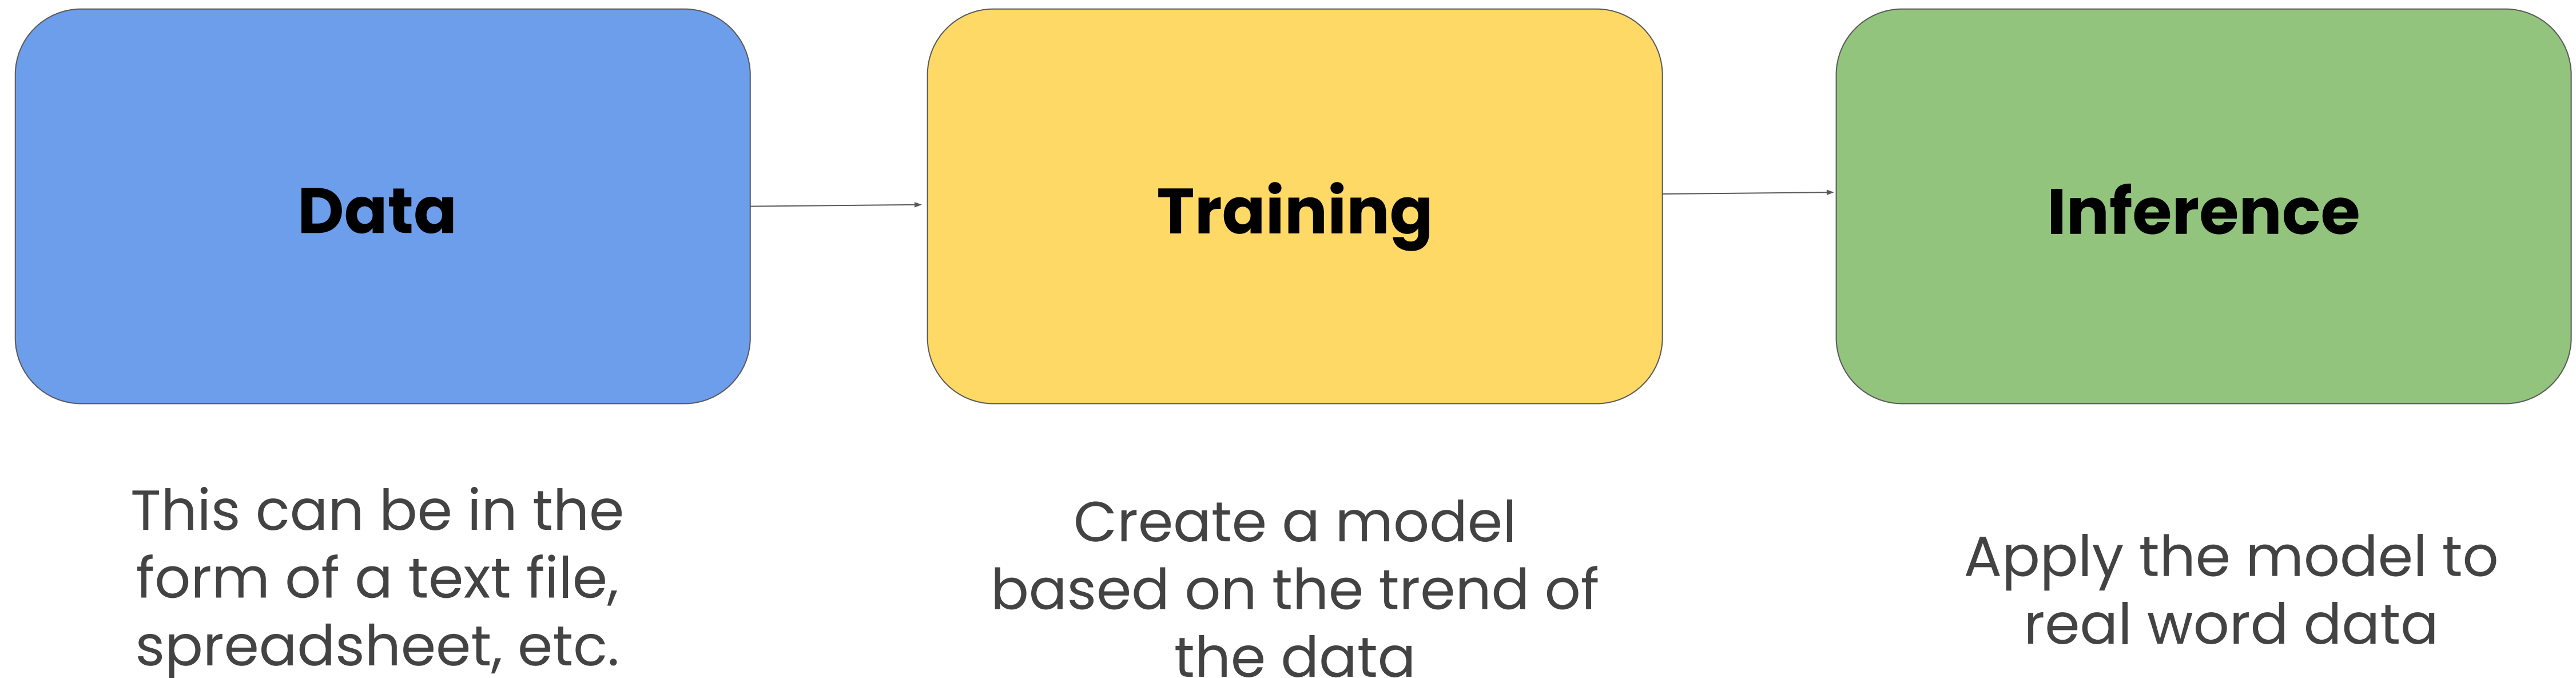
Eg. Takes in a sentence and translates it into French

- So it sounds like it would be pretty nice to have a model.
Machine Learning helps us create such models through a process called *training*

Now that we know more about models...

- How would we predict dorm pricing with a model?
- What are some possible inputs for our model?
 - Think about what you would need to tell how expensive a dorm room is
- What would the output of our model look like?
 - Would it be continuous, or would it be categorical (“this or that”)?
- How would we go about *training* such a model?

ML Pipeline



More on the intuition

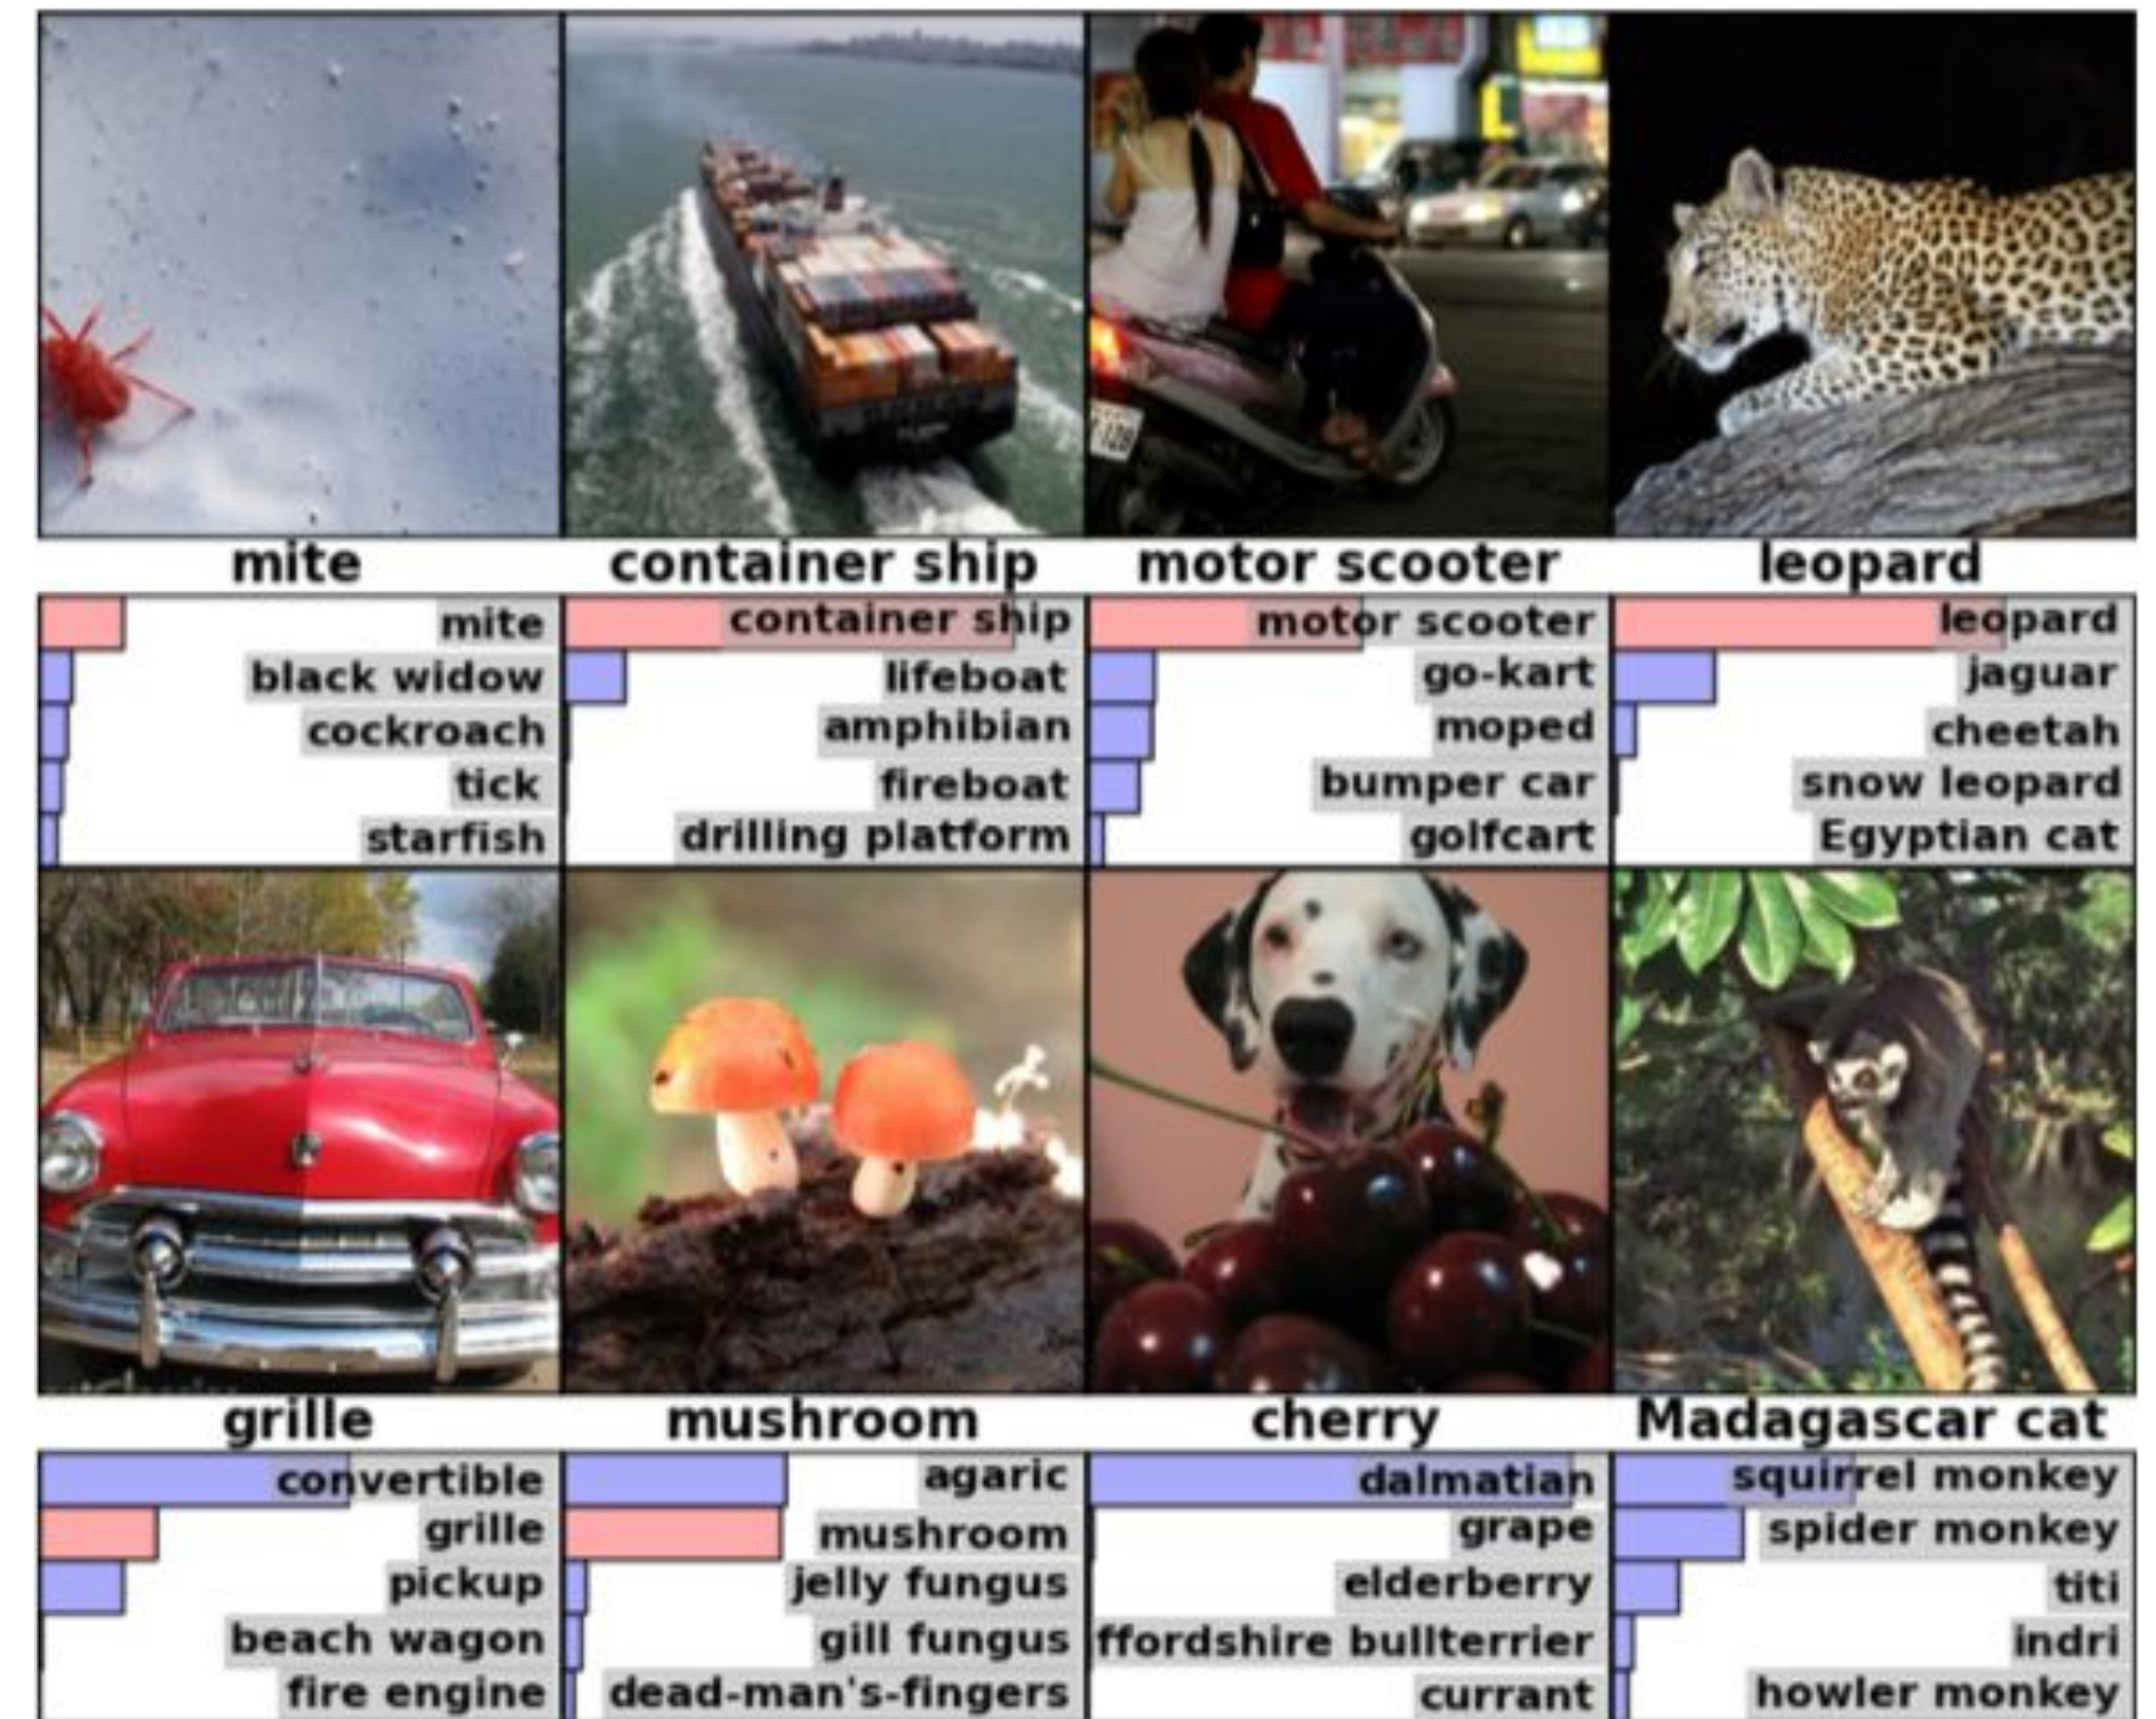
How do you know the difference between a cat and a dog?

- Did someone teach you what to look for?
- What specific features distinguish them?
- Every time you got it wrong, your parents told you what animal it was
- Eventually, you learned how to tell the difference
- This is how we will eventually improve our model as well!



Recognizing Objects in Images: AlexNet

- Trained on millions of photos of different objects
- Learned to **classify** different images by slowly recognizing patterns
- [Paper](#)



Python + Environment setup

- We will be using **Google Colab** notebooks, which will come with all the packages pre-installed.
- The **Anaconda Distribution** is *not required* for this workshop series, but it's a great tool to work with Jupyter notebooks in general.

Intro to Python



- We're going to now take a 5 minute break
- The next part of this workshop is going to be a quick introduction to **Python** for those of you who haven't used it before
- If you do feel comfortable with Python already, feel free to leave and come back for our next workshop when we start talking about some ML models!

Take a break for 5 minutes!

Link to the Python Tutorial

<https://tinyurl.com/btrack-w22-python>

Thank you! We'll see you next week!

Please fill out our feedback form:

bit.ly/btrack-w22-feedback

Next week: **The K-Nearest-Neighbours Model**

FB group: facebook.com/groups/uclaacmai

